


Style Transfer

By Rediet Negash


Have you ever
dreamed of
painting like Van
Gogh?



Well, even if u don't have the skills to do so, don't be sad

you still can do this with a super cool subject called Style transfer, in which the style of a piece of artwork is transferred onto a picture.

Eg: SUNY Korea + Starry Night + other style






Styles used can be...




How does this work?

Style transfer is an optimization problem where the neural network is not required to do something but rather the backpropagation and optimization was used to slowly alter the image to incorporate style characteristics.



Neural Style



$$\mathcal{L}_{total}(\vec{p}, \vec{a}, \vec{x}) = \alpha \mathcal{L}_{content}(\vec{p}, \vec{x}) + \beta \mathcal{L}_{style}(\vec{a}, \vec{x})$$

Match activations
from content image

Match gram matrices
from style image

Content loss:

$$\mathcal{L}_{content} = \sum_l \sum_{i,j} (\alpha C_{i,j}^l - \alpha P_{i,j}^l)^2,$$

Style Loss

Similar with content loss computation except instead of comparing the raw outputs of the style and the target images at various layers, we compare the **Gram matrices** of the outputs.

A Gram matrix results from multiplying a matrix with the transpose of itself:

$$G_{i,k}^l = \sum_k F_{i,k}^l F_{j,k}^l$$

Why Gram Matrix



Since every column is multiplied with every row in the matrix, the spatial information that was contained in the original representations could be considered to have been “distributed”. The Gram matrix instead contains non-localized information about the image, such as texture, shapes, and weights - which is what we want for style transfer

$$\mathcal{L}_{style} = \sum_l \sum_{i,j} (\beta G_{i,j}^{s,l} - \beta G_{i,j}^{p,l})^2.$$

Experiments

—

Objective

- Improve aesthetics of a target Image
- Multiple styles
- Photo style transfer
- Improve Denoising resulted



Approaches used

Hyper parameter tuning

1, Doing different number of Iterations that results in a desired image

2, Changing the convolutional layers used

- For the style layers
- For content layers

3, Changing the weights of the style and content losses (alpha and beta)

4, changing the optimizers hyper parameters- learning_rate

Frameworks used

Torch

VGG19

Torchvision:-

Python Imaging Library(PIL).

For content comparison I used activation/feature map at some deeper layer conv4_2 and conv4_2 layers prior to the output(softmax) layer.

Adam Optimizer

```
[36]: import torch
      from PIL import Image
      from torchvision import transforms, models
      import numpy as np
      import matplotlib.pyplot as plt
```

```

In [37]: model = models.vgg19(pretrained=True).features

# freeze the parameter that we don't use to minimize computation and memory
for parameter in model.parameters():
    parameter.requires_grad = False

#print(model)

# extract features
def extract_features(image, model):
    features = {}
    layers = {
        '0' : 'conv1_1',
        '5' : 'conv2_1',
        '10' : 'conv3_1',
        '19' : 'conv4_1',
        '21' : 'conv4_2', # for measuring content loss
        '28' : 'conv5_1'
    }
    feature_image = image
    # apply singleton dimension image
    feature_image = feature_image.unsqueeze(0)
    # pass the image to layers in the model
    for key, layer in model._modules.items():
        feature_image = layer(feature_image)
        # if the current layer in the model is part of the layers
        # planned to use add activated feature to the features dictionary
        if key in layers:
            features[layers[key]] = feature_image
            #print(feature_image)

    return features

```

VGG model

Model layers

Sequential(

(0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(1): ReLU(inplace=True)

(2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(3): ReLU(inplace=True)

(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

(5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(6): ReLU(inplace=True)

(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

(8): ReLU(inplace=True)...

model._modules.items()

[('0', Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))), ('1', ReLU(inplace=True)), ('2', Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))), ('3', ReLU(inplace=True)), ('4', MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)), ('5', Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))), ('6', ReLU(inplace=True)), ('7', Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))), ('8', ReLU(inplace=True)), ('9', MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)), ('10', Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))), ('11', ReLU(inplace=True)),...

Helper functions used

```
transform = transforms.Compose([transforms.Resize(300),
                                transforms.ToTensor(),
                                transforms.Normalize((0.5,0.5, 0.5),(0.5,0.5,0.5))])

def tensor_to_Img(tensor):
    img = tensor.clone().detach().numpy().squeeze()
    img = img.transpose(1,2,0)
    img = img*np.array((0.5,0.5,0.5)) + np.array((0.5,0.5,0.5))
    return img
def img_show(content, style):
    fig, (ax1,ax2) = plt.subplots(1,2)
    ax1.imshow(tensor_to_Img(content), label = "Content")
    ax2.imshow(tensor_to_Img(style), label = "Style")
    plt.show()
```

```

def create_style_loss(style, output_features):

    style_features = extract_features(style,model)
    style_weight = {"conv1_1" : 0.4,
                    "conv2_1" : 0.3,
                    "conv3_1" : 0.2,
                    "conv4_1" : 0.2,
                    "conv5_1" : 0.2

                    "conv1_1" : 1.0,
                    "conv2_1" : 0.8,
                    "conv3_1" : 0.4,
                    "conv4_1" : 0.2,
                    "conv5_1" : 0.1
                    }

    # compute the correlation (gram matrix)
    gram_layers = {}
    for layer in style_features:
        gram_layers[layer]= gram_matrix(style_features[layer])

    # calculate style loss
    style_loss = 0
    for layer in style_weight:
        style_gram = gram_layers[layer]

        output_gram = output_features[layer]
        # used for normalization
        _,depth, weight, height = output_gram.shape
        output_gram = gram_matrix(output_gram)
        # squared mean of the output and the style gram
        style_loss += (style_weight[layer]*torch.mean((output_gram-style_gram)**2))/depth*weight*height

    return style_loss

```

```
def gram_matrix(feature):
    t, depth, height, weight = feature.size()
    #input = torch.Tensor(2, 4, 3) # input: 2 x 4 x 3
    #print(input.view(1, -1, -1, -1).size()) # prints - torch.size([1, 2, 4, 3])

    feature = feature.view(depth, height*weight)
    # multiply the matrix with its transpose
    gram_matrix = torch.mm(feature, feature.t())

    return gram_matrix

def create_content_loss(content, output_features):

    content_features = extract_features(content,model)
    content_loss = torch.mean((content_features['conv4_2']-output_features['conv4_2'])**2)
    return content_loss
```

```

def main():
    # alpha
    content_wt = 5
    # beta
    style_wt = 1e9
    # style_wt2 = 1e6
    # style_wt3 = 1e6

    print_after = 100
    epochs = 2000

    content = Image.open("img_red.jpg").convert("RGB")
    content = transform(content)

    style = Image.open("1-style.jpg").convert("RGB")
    style = transform(style)

    # style2 = Image.open("2-style1.jpg").convert("RGB")
    # style2 = transform(style2)

    # style3 = Image.open("style.jpg").convert("RGB")
    # style3 = transform(style3)

    img_show(content, style)
    # img_show(style3, style2)

    # output is initialized from the content first
    output = content.clone().requires_grad_(True)

    # optimize the output image
    # this is a hyper parameter
    optimizer = torch.optim.Adam([output], lr=0.001)
    # lr - learning rate, 0.07
    # activate output features
    for i in range(1, epochs + 1):
        output_features = extract_features(output, model)

        content_loss = create_content_loss(content, output_features)
        style_loss = create_style_loss(style, output_features)
        # style_loss2 = create_style_loss(style2, output_features)
        # style_loss3 = create_style_loss(style3, output_features)

```

```

        # activate output features
        for i in range(1, epochs + 1):
            output_features = extract_features(output, model)

            content_loss = create_content_loss(content, output_features)
            style_loss = create_style_loss(style, output_features)
            # style_loss2 = create_style_loss(style2, output_features)
            # style_loss3 = create_style_loss(style3, output_features)

            total_loss = content_wt*content_loss + style_wt*style_loss
            # total_loss = content_wt*content_loss + style_wt*style_loss + st

            if i % 10 == 0:
                print("epoch ", i, " ", total_loss)

            optimizer.zero_grad()
            total_loss.backward() # This calculates the gradients
            optimizer.step() # This updates the net

            if i % print_after == 0:
                plt.imshow(tensor_to_img(output), label="Epoch "+str(i))
                plt.show()
                # plt.imsave(str(i)+'.png', tensor_to_img(output), format='png')

if __name__ == '__main__':
    main()

```





```

if i % print_after == 0:
    plt.imshow(tensor_to_img(output), label="Epoch "+str(i))
    plt.show()
    # plt.imsave(str(i)+'.png', tensor_to_img(output), format='png')

if __name__ == '__main__':
    main()



```



```



epoch 10  tensor(4.7925e+19, grad_fn=<AddBackward0>)
epoch 20  tensor(4.3304e+19, grad_fn=<AddBackward0>)
epoch 30  tensor(3.8863e+19, grad_fn=<AddBackward0>)
epoch 40  tensor(3.4605e+19, grad_fn=<AddBackward0>)

```





Some outputs

—



More Iterations with multiple styles








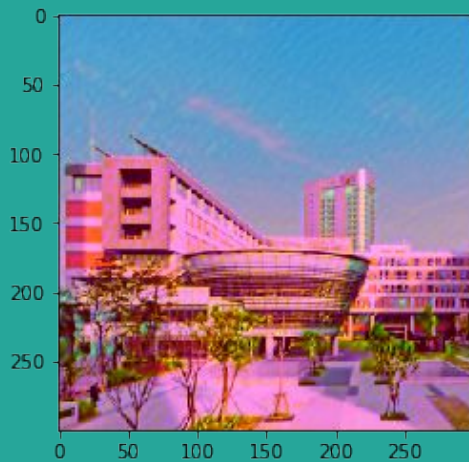
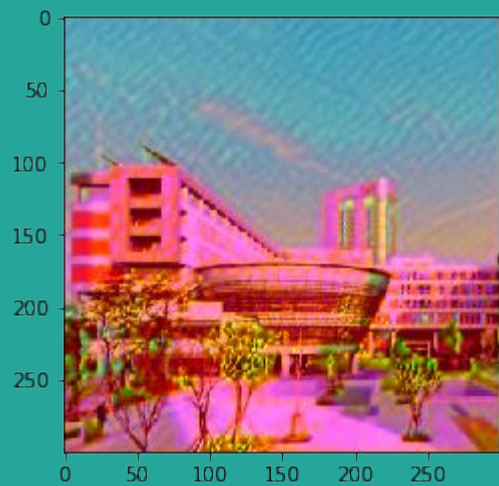
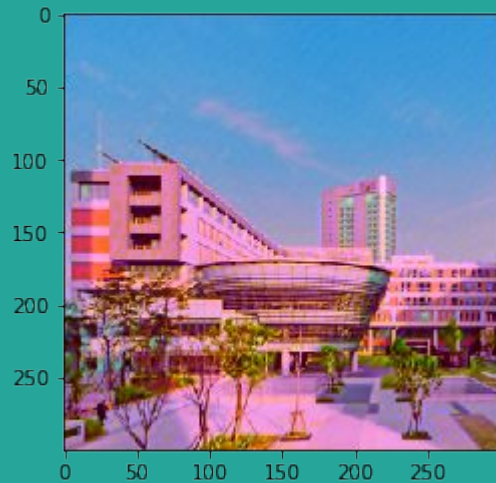




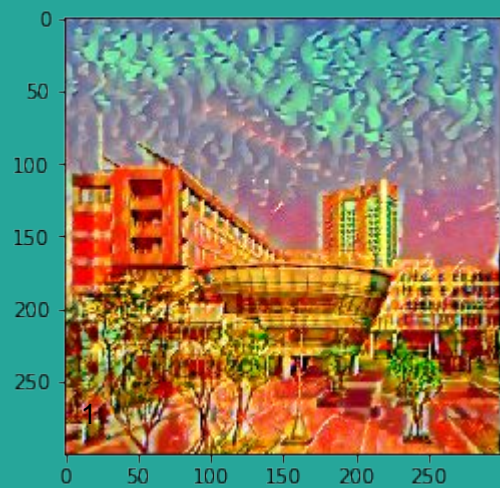




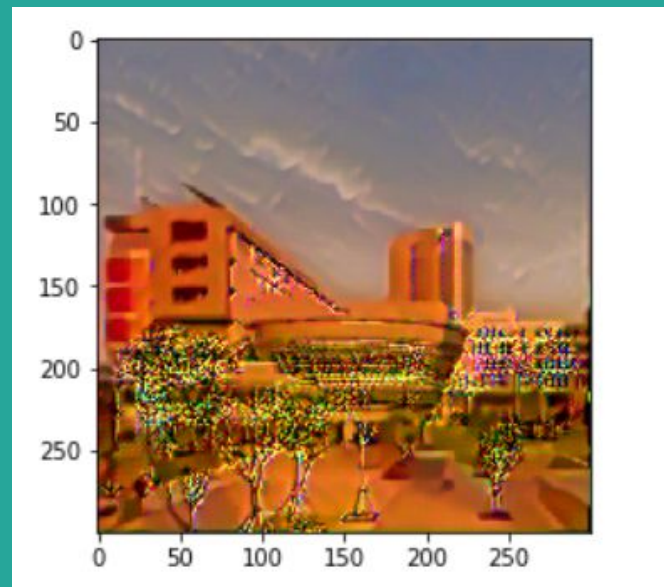

Photo realistic style transfer

—





Two Styles applied



Future and current works

- **Feedforward Method**
- **Arbitrary Style Transfer**
- **Improve computational time**
- **Improve aesthetics of generated image_F**

Conclusion

- The quality of a style transferred image highly depends on the hyper parameters and the convolutional layers used.
- Using higher layers for content and lower layers for style does give a good quality of a generated image. Again, quality is subjective.

Multiple styles- Getting a good neural styled image does require an artistic mind where the content and the style image should be carefully chosen to get the best results.

*** future work teaching the algorithm to predict best style combination for synthesising the best combination of styles...

Thank you

Presented by Rediet Negash

